


- git reset: undo add
- git revert: undo commit
- Each git repository can act as a server (provider of commits) or client
- git objects
  - Commit: a reference to a change being tracked by git
    - Refers to a tree object that describes the state of the repo for that commit
  - Tree: A directory being tracked
    - Contains a listing of blobs and tree objects (subdirectories)
  - Blob: a file being tracked by git
- refs
  - heads/: a listing of commit-type object DB keys that each branch currently points to
  - remotes/: a list of branches that are currently being tracked in upstream repositories
  - tags: A list of commit-type object DB keys that each tag corresponds to
- HEAD: a reference to the current branch that is checked out
- config: a file that stores config options for the current repo
- pre-receive: takes a list of references that are being pushed from stdin
  - e.g. make sure none of the updated refs are non-fast-forwards
- update: run once for each branch being pushed
  - name of ref (branch), the SHA-1 that ref pointed to before the push, and the SHA-1 that is being pushed
  - e.g. reject pushes to master
- post-receive: runs after the entire process is completed; takes same stdin input
  - e.g. notify a CI server or email a list

- Build tool families:
  - Low-level:
    - Explicitly define deps and rules for each input and output file
  - Abstraction-based:
    - Derive low-level build code from high-level data, e.g. maps of files to executables
  - Framework-driven
    - Default behavior assumed unless explicitly overridden
- Complementary tool families
  - Dependency management
  - Testing frameworks
- Makefile variables
  - `my-var = "hello"`
  - `random.o: random.c`
  - `<tab> $(my-var)`
- Incremental builds: only out-of-date targets are rebuilt
- Low-level build issues
  - Shell scripts aren't super portable
- To handle deps across directories, recursive calls to Make are typically used
  - Fractures global dependency graph
  - One build may not rebuild all that is needed
- Ant is one low-level solution
  - Properties (variables)
  - Tasks (invocation of commands)
  - Targets (grouping of tasks)
- Maven:
  - A lifecycle is a sequential series of phases
  - A phase performs a series of goals that are bound via plugins
  - Default, Site, Clean
- A phase with no goals or plugins is skipped
- Dependencies can be scoped to specific phases
- Builds are computationally expensive
  - Dependency graph requires a lot of RAM
  - Compilation requires lots of CPU
  - File staging requires large and fast disks
- Builds can easily be parallelized, but there is a limit on developer machines
- Bazel creates an action graph (BUILD & WORKSPACE)





- Devs execute personal builds on their machines
- Only quick tests are run by default
  - Slower tests can be relegated to special build targets that run less often
- **Micro-build:** Concerns about the behaviour of a build system within a single execution
- **Macro-build:** Concerns about how to best provision a fleet of resources
- Teams that use modern code review connect build jobs to the code reviewing dashboard
- Nightly builds are too infrequent
  - If hundreds of devs made commits, it's hard to tell who caused the problem
  - Imagine your code from yesterday broke the build — tough to recall what you were doing

### • **CI Feedback loop:**

- Commit, Build, Test, Report
- The benefits of CI
  - Each commit gets its own build job
  - Errors are reported quickly
- Build systems need to be updated as the codebase evolves
  - Otherwise build can produce incorrect results
  - Broken builds + weird bugs

### • **Flaky tests** (non-deterministic)

- False positive: test says code has failed but it should have passed
  - Devs lose trust in the test
- False negative: test passes when it should have failed
  - Allows bugs to slip through
- The Modern CI Process
  - Build - triggering event
    - Triggered by push or manually by dev
  - Build job creation service
    - Build job creation node adds the job to a queue of pending jobs
  - Build job processing service
    - Build jobs in the pending queue are allocated to build job processing nodes
    - First download latest version of code then apply changes to it

- Compile, execute automated tests, deploy
- Add results to report queue
- Build reporting service
  - Results communicated to dev team
- Advantages of inspections vs. dynamic QA
  - Cascading errors can obfuscate test results
  - Incomplete work can still be inspected
  - Can uncover inefficiencies and style issues
  - Knowledge sharing, collaborative problem solving
- Inspections complement tests
- Structured inspections (rigid, heavyweight)
  - Moderator ensures...
    - Artifact is ready for review
    - Inspection procedure is followed
  - Scribe/Recorder
  - Reviewer
  - Reader (leads team through the code)
  - Producer (author of artifact)

### • Tasks in software inspection

- Planning (find review team, time, place)
- Group prep, individual prep
- Checklists ensure quality is attained
- Tips for productive review

#### DO:

- Critique the artifact
- Keep review chunks short and succinct
- Plan time for reviews
- Prioritize reviews of important issues
- Keep it light

#### DON'T:

- Attack the person
- Submit fixes for multiple issues at once
- Skip reviews
- FIFO queue reviews
- Use sarcasm or exaggeration

### • Why do people do code reviews?

- Finding defects
- Code improvement
- Team awareness
- Share code ownership
- Alternative solutions
- Knowledge transfer
- Improve dev process
- Avoid build breaks

### • Modern review roles

- Author
- Reviewer
- Integrator (makes final call)
- Verifier (checks the functional correctness of the code)
- Some Maven phases:
  - test
  - clean
  - compile
  - test-compile
  - deploy
  - install



What does virtualization provide?

## Ways to find defects

- Testing
- Static & Dynamic Analysis Tools
- Statistical Defect Prediction
- Manual Inspections

## Automatic defect finding

- Goal: discover likely locations of defects

- Static analysis: code smells, bad patterns, definite errors

- Bug prediction: ML & Statistics

- Dynamic ~~analysis~~ <sup>analysis</sup>: taint analysis, performance, memory

## Bug localization (Hybrid Dynamic/Stats)

- Use statistical properties to calculate suspiciousness of a line location

## Code Smells

- Duplicated code
- Long methods
- God classes
- Long parameter lists
- Primitive obsession

## Static Analysis

- Syntax Analysis
- Semantic (data flow analysis)

## Data Flow Analysis

- Framework for proving facts about a program
- Examines how information propagates through blocks and paths of a program

## Reaching Definitions

- A definition of a variable  $x$  is a statement that may modify its value
- A use of a variable  $x$  is a statement that reads from  $x$

• DU-chain: links each def to uses it reaches

• UD-chain: links each use to reaching defs

## Virtual Machines (VMs)

- Emulate a computer system
- Provide a computer system from within the scope of another system

- Containers
  - ~~are~~ A set of processes running on top of a common kernel
  - Isolated from each other

- Makes entire machines "shippable" (OS, system hardware application stack)
- The need for virtualization:

- Gives shared infrastructure users (e.g. cloud users) the impression of having their own machine

## System virtualization:


- "Full" virtualization
- Hosts an entire OS

## Process Virtualization


- Provides the isolation of system virtualization while also sharing common software tools with the underlying platform

• Hypervisor: a computer software, firmware or hardware that creates and runs virtual machines (aka Virtual Machine Monitor)

## Type 1 Hypervisor:



## Type 2 Hypervisor:



## Namespaces

- Enable processes to have a private view of the system resources
- Network interfaces, Process IDs, Mount points

## Docker

- An orchestration tool for containers. Features:
- Portability, App-centric, Builds from source, versioning, Component reuse, public registry, tool ecosystem

## Before containers:

- Very inefficient use of memory and CPU resources

## After containers:


- Isolated services in fewer VMs
- Uses VMs more efficiently

## Docker Engine has:


- A server (daemon)
- A REST API which specifies interfaces that programs can use to talk to the daemon
- A CLI client (the "docker" command)

## Blue-Green Schema changes

- Transitional status, DB schema should work for both new & old layers
- Final cleanup: delete old schema data



# Docker Engine Visualized:



• Docker daemon: listens for Docker API requests and manages images, containers, networks and volumes

• Docker Registry: stores Docker images (i.e. Docker Hub)

## The Cloud:

- Shared pools of configurable computing resources
- Can be rapidly provisioned (semi)-automatically
- Can "elastically" scale up or down based on demand
- Only pay for what you use

## Downsides of cloud computing

- Security, privacy, and other concerns about data and IP shared in the cloud are real

## IaaS

- Provides computing resources to users in the form of VMs or containers
- Compute hours/months or per-machine sizing (small, medium, large)

## IaaS vs. In-house Infrastructure

- + Explainable infra operation costs
- + IT support offloaded to cloud provider
- + Elastic scaling
- Data security & privacy may be diff. to verify
- Costs can grow in unexpected ways

## PaaS

- Service that provides a development platform
- Provides users with comfortable app development entry point

## PaaS vs. IaaS

- + Devs can focus on app development instead of config
- + Platforms are often tuned to near-optimal solutions

- Convenience at the cost of configurability
- Non-standard workloads may suffer
- Hard to port applications to other platforms

## SaaS

- The entire stack (infra, platform, execution env) is hosted on the cloud and delivered over the web
- Usually involves a "thin client" that connects to a backend that delivers most of the functionality
- e.g. Google suite

## Serverless

- Variant of IaaS where capacity planning decisions are made by the provider (instead of consumer)
- Customers are charged for the resources that their apps use
- Elastic scalability is at app level instead of VM level

## Serverless vs. IaaS

- + Users pay only for what they use
- Costs can be tough to predict

## FaaS

- Primary means by which serverless computing is achieved
- Functions are deployed to be executed on-demand in the cloud

## MBaaS

- Share common backend features as APIs
- e.g. Authorization, notif routing, social media integration, cloud-like storage
- e.g. Firebase

## Infrastructure Config

- Difficult to do manually
- REs spend lots of time firefighting with config details
- Wastes precious person-hours

## IaC

- Write configurations as code
- Manage and provision machines using a code-like syntax rather than interactive config tools
- Has all advantages of code
- Versioned (VCS)
- Can be automatically executed

## Declarative Programming

- Expresses the logic of a computation w/o describing its control flow
- Describes what the program needs to do, and not how



- Partial time-limited deployment
- Enforce safety of changes

## Docker commands

- docker build
- docker run
- docker start/stop
- Dockerfile commands
- FROM, COPY, RUN, CMD, WORKDIR
- Puppet
- Resource: lets you inspect & manipulate resources on a system
- ldd: prints shared libraries that an executable depends on

```
class pasture {
  package { 'pasture':
    ensure => present,
    provider => 'gems',
  }
  file { ['/path':
    source => puppet:///...
  ]
  service { 'pasture':
    ensure => running,
  }
}
```

$$P = \frac{(a+b)(c+d)}{(a+c)(b+d)}$$



- Roll forward
- Roll backward
- Alert certain parties